Visual Question Answering

..because it caught our attention

Motivation

1. Visually impaired users
2. Information retrieval for a chain of tasks

What color are her eyes?
What is the mustache made of?

How many slices of pizza are there?
Is this a vegetarian pizza?

Is this person expecting company?
What is just under the tree?

Does it appear to be rainy?
Does this person have 20/20 vision?

Dataset and Evaluation Metrics

1. Images — COCO dataset
2. Questions — 3 unique individuals
3. Answers — 10 unique individuals
4. Accuracy = Minimum (correct answers/3, 1)

Baseline

1. Image Encoding: VGG19
2. Question Encoding: LSTM
3. Combined the inputs using element wise multiplication
4. Passed through two fully connected layers before finally taking a softmax

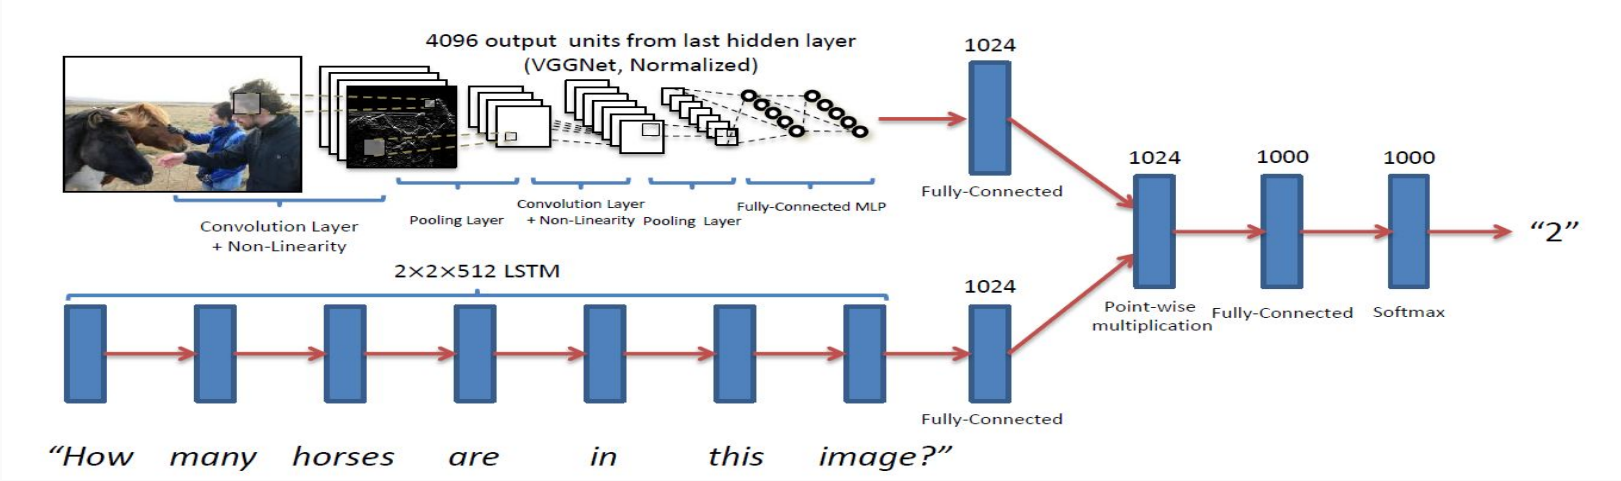

Attention

1. Attention models extract more information
2. Higher weights are put on the visual regions that are more relevant to the question.

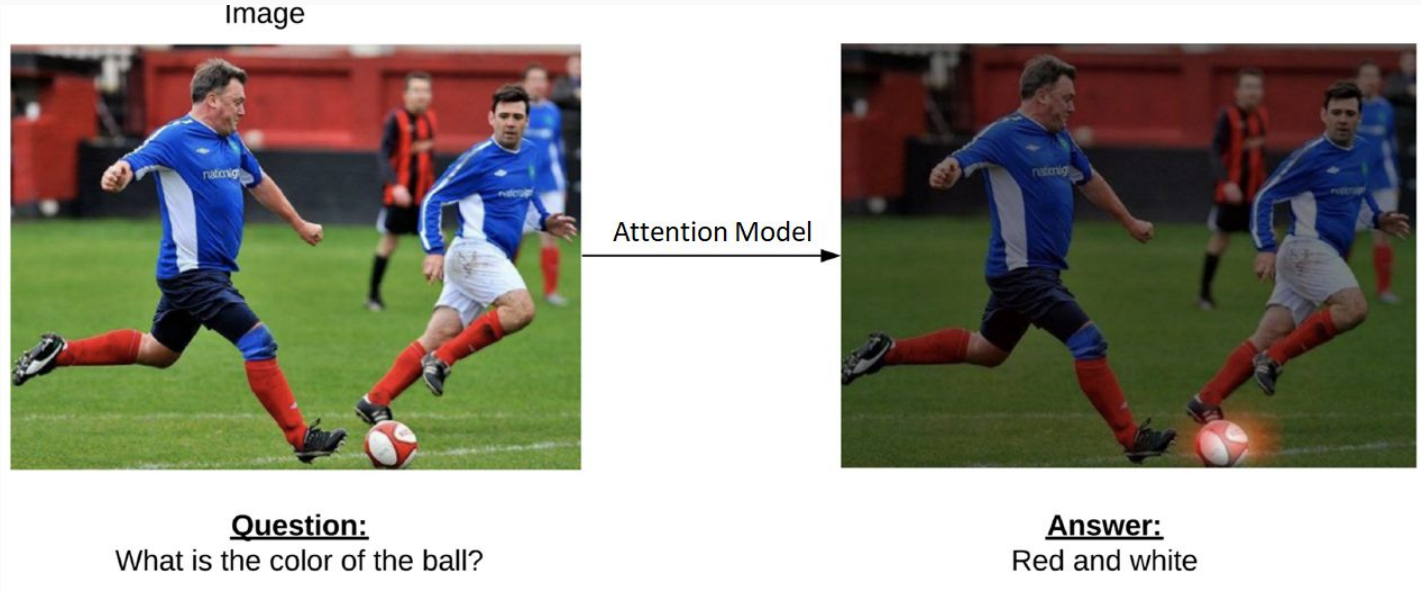

Stacked Attention

1. Iterate the above query-attention process using multiple attention layers
2. Extracts more fine-grained visual attention information for answer prediction

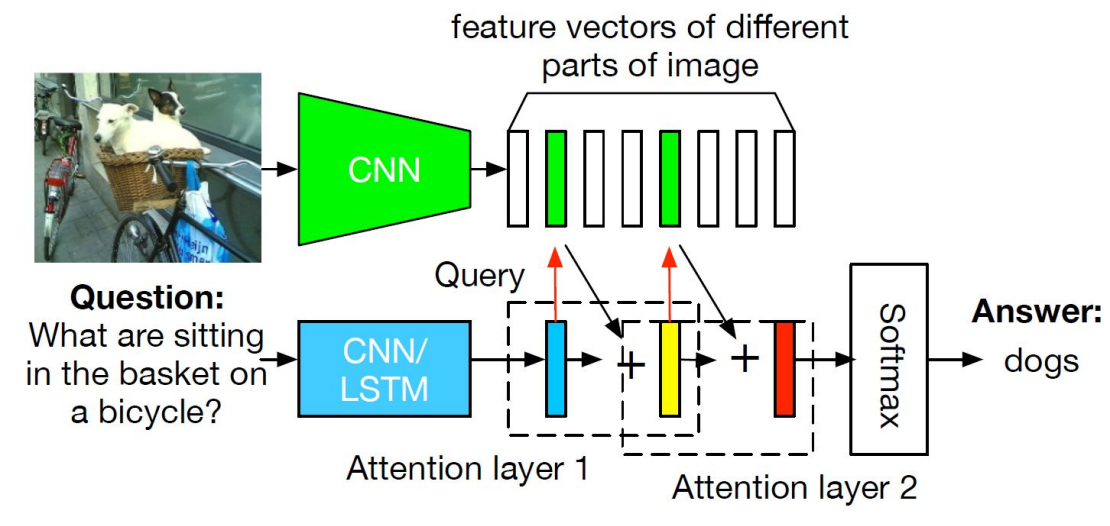

Visualizing Stacked Attention

Question: What are sitting in the basket on a bicycle?

Bottom-Up and Top -Down Attention

- All the previous approaches on attention can be summarized as top-down approaches
- They make the model focus on particular region of the image based on the state of the current task (e.g. looking for something).
- However, in this approach all the image features are divided into uniform grid which are then weighted by the top-down attention.
- Dividing image into regions by interesting objects is a better idea.

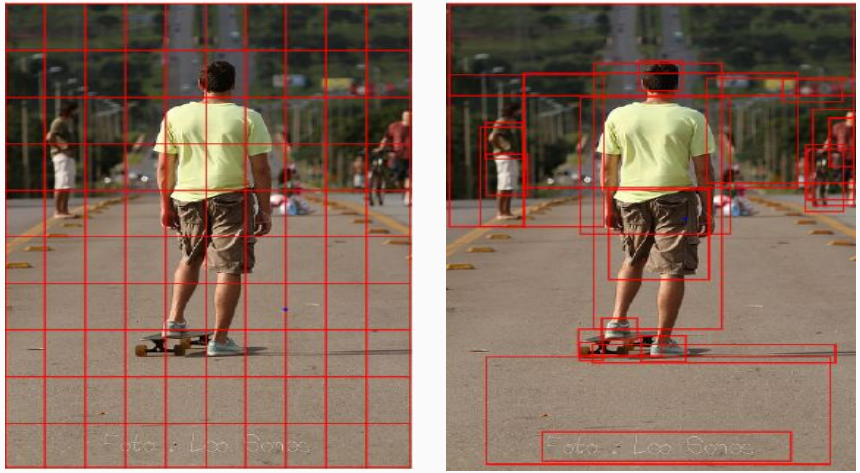

In the left image the image regions are divided into uniform grid which is weighted by attention. In the right,weight image by interesting regions.

Architecture

- Hence, the authors propose a bottom-up attention model. This visual encoder provides the encoded image features as well as the objects of interest.
- This is achieved by using a **Faster R-CNN**.
- The first stage predicts object proposals and the second stage extracts a feature map for each box proposal.

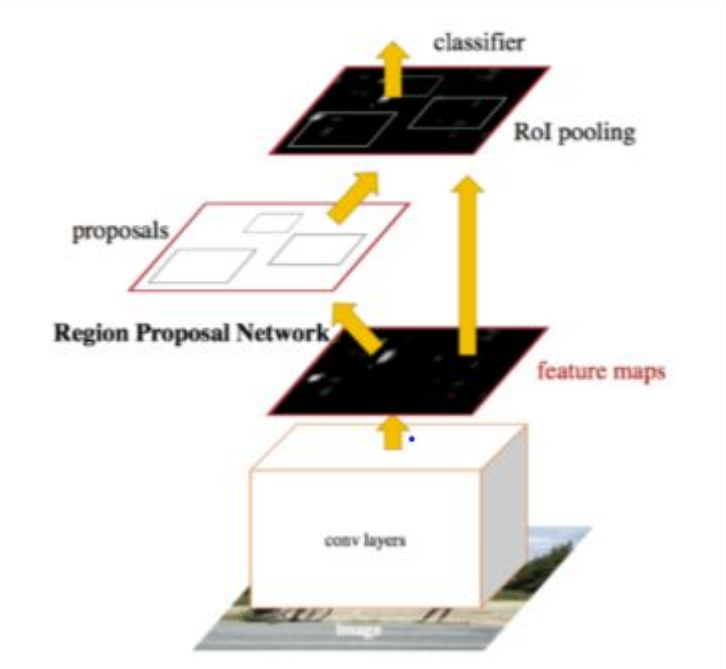

Results

Results on Validation Data Set

Model	**Validation Accuracy**
Baseline	**57.50%**
Simple Attention	**59.01%**
Two Layered Stacked Attention	**59.98%**
Four Layered Stacked Attention	**60.48%**
Bottom-up and Top-Down Attention	**62.71%**

Demo Images

Question: what sport is being played? Correct Answer: baseball

\Model	Baseline	Simple Attention	Stacked Attention (2)	Stacked Attention (4)
Prediction 1	tennis (0.34)	**baseball (0.20)**	**baseball (0.26)**	**baseball (0.28)**
Prediction 2	soccer (0.07)	tennis (0.10)	tennis (0.17)	tennis (0.25)
Prediction 3	unk (0.04)	snowboarding (0.06)	snowboarding(0.16)	skateboarding (0.09)
Prediction 4	skateboarding (0.04)	unk (0.03)	skateboarding (0.05)	basketball (0.06)
Prediction 5	**baseball (0.03)**	none (0.03)	skiing (0.04)	frisbee (0.04)

Question: is there a train? Correct Answer: yes

Prediction\Model	Baseline	Simple Attention	Stacked Attention (2)	Stacked Attention (4)
Prediction 1	**yes (0.63)**	**yes (0.66)**	**yes (0.55)**	**yes (0.50)**
Prediction 2	no (0.36)	no (0.31)	no (0.41)	no (0.41)
Prediction 3	maybe (0.001)	maybe (0.002)	window (0.0037)	maybe (0.0064)
Prediction 4	parked (0.000)	unknown(0.002)	reflection (0.0023)	black and white (0.0037)
Prediction 5	safety (0.000)	1 (0.0015)	maybe (0.0019)	maybe (0.0037)

Question: how many people are there? Correct Answer: 4

Prediction\Model	Baseline	Simple Attention	Stacked Attention (2)	Stacked Attention (4)
Prediction 1	2 (0.23)	1 (0.22)	1 (0.10)	1 (0.19)
Prediction 2	1 (0.22)	2 (0.21)	2 (0.07)	2 (0.11)
Prediction 3	3 (0.14)	3 (0.06)	0 (0.06)	**4 (0.08)**
Prediction 4	**4 (0.10)**	**4 (0.06)**	3 (0.05)	0 (0.08)
Prediction 5	0 (0.08)	many (0.04)	10 (0.05)	3 (0.06)

Next Steps

1. Combine lower, middle, and final layers of CNN
2. Bilinear Attention Models
3. Mask R-CNNs

Summary

1. Various models for VQA
2. Improved performance with the use of attention
3. Plausible answers obtained on new images

Questions?

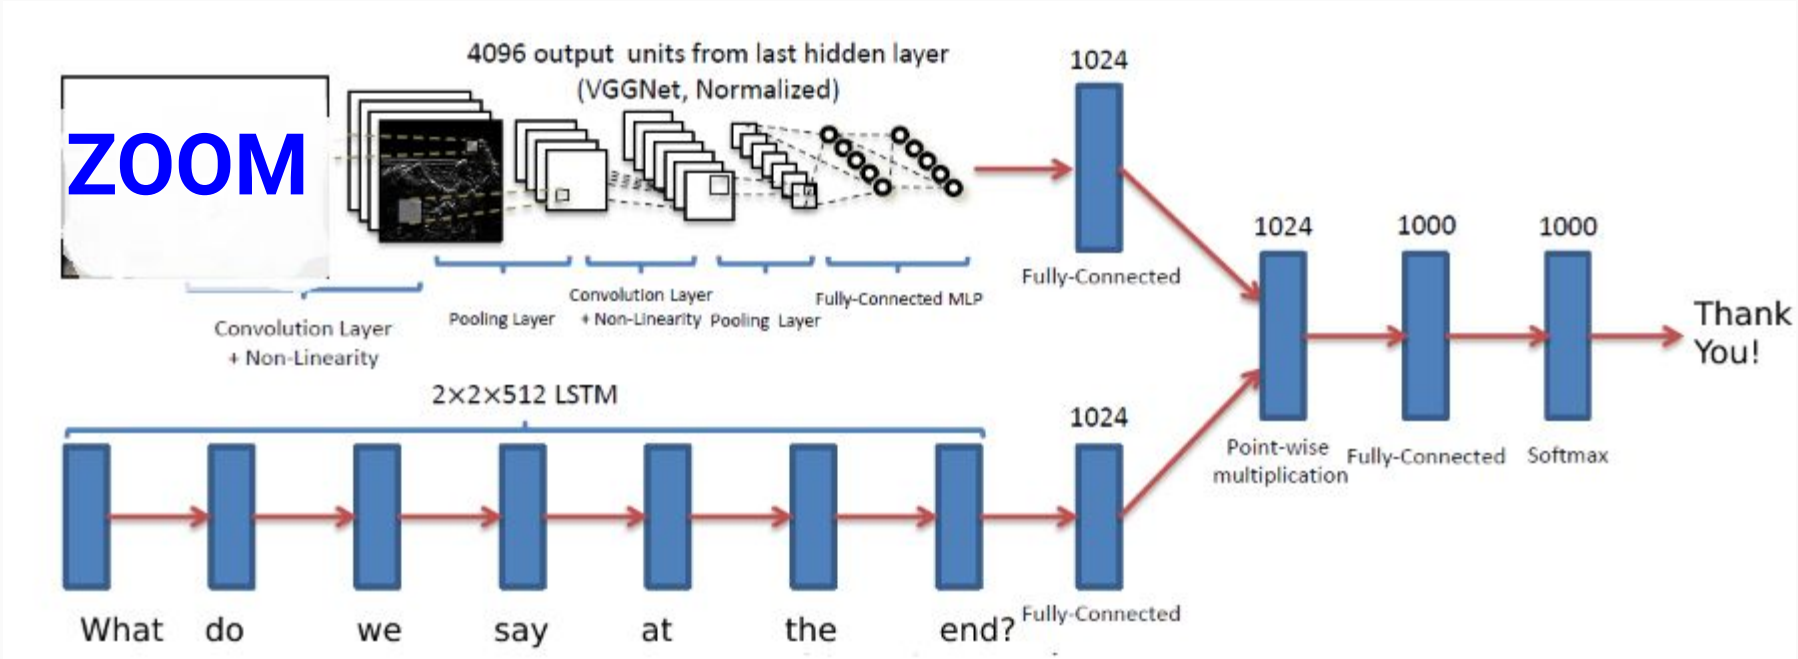
ZOOM
4096 output units from last hidden layer
(VGGNet, Normalized)
1024
Fully-Connected
Convolution Layer
+ Non-Linearity
Pooling Layer
Convolution Layer
+ Non-Linearity
Pooling Layer
Fully-Connected MLP
2×2×512 LSTM
1024
1024
1000
1000
Point-wise multiplication
Fully-Connected
Softmax
Thank You!
What
do
we
say
at
the
end?
Fully-Connected